A4 learned residual scales

Screening step 1276 | seed 42 | A0 fixed scale = 0.1

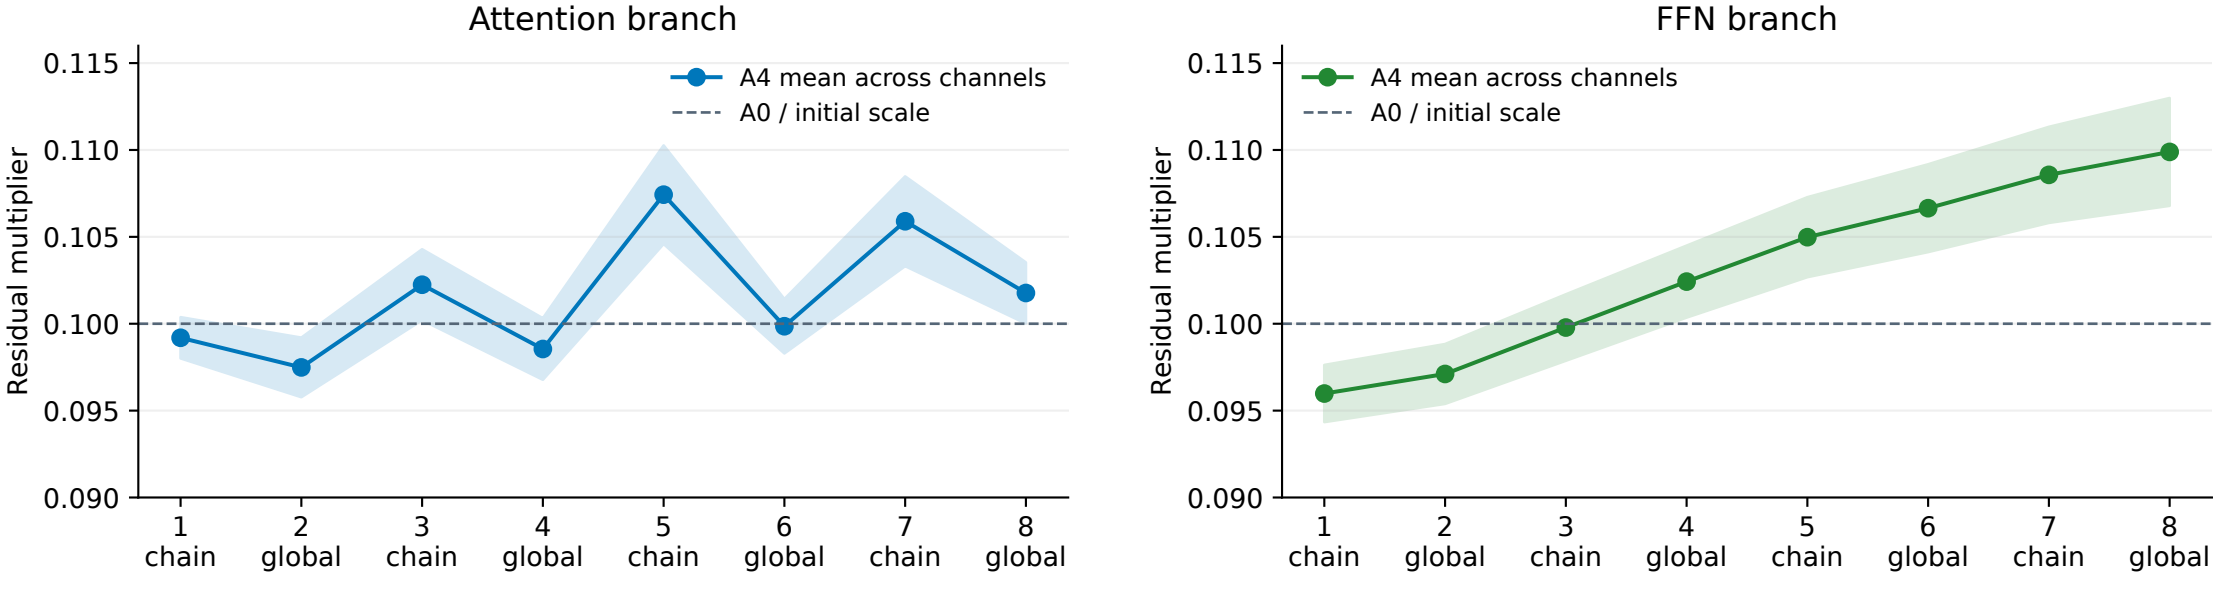

Layer number and its attention visibility. Shading: channel SD; this is not variation across training seeds.

Descriptive weight analysis only: other weights adapt jointly, so scale magnitudes do not establish causal feature contributions.