Architecture screening | completed-run comparison

2026-10-28 22:54 UTC+8 | A0, A1, A2, A3, A4, A5, A6: final step 1276 | same 10% V8 subset | seed 42

A0 Baseline A1 QK norm A2 Attention gate A3 QK norm + gate
A4 Learnable residual scale A5 All-global attention A6 Tapered FFN

Whole-chain generation: AAR

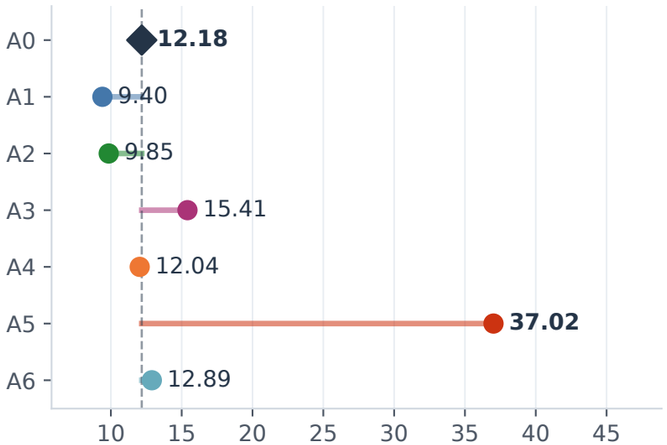

Higher is better | %

Whole-chain generation: length MAE

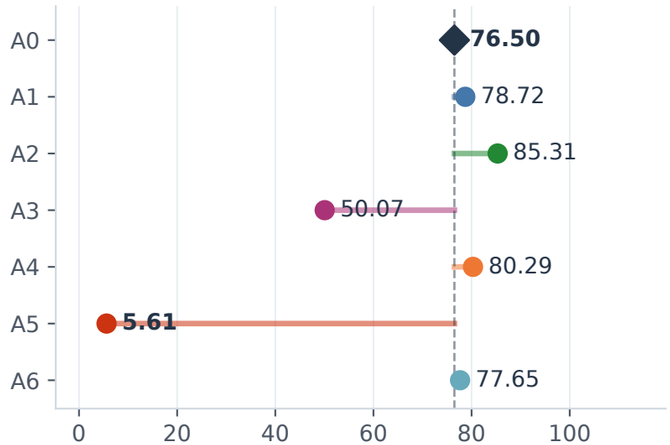

Lower is better | residues; valid outputs only

Representation: macro AUROC

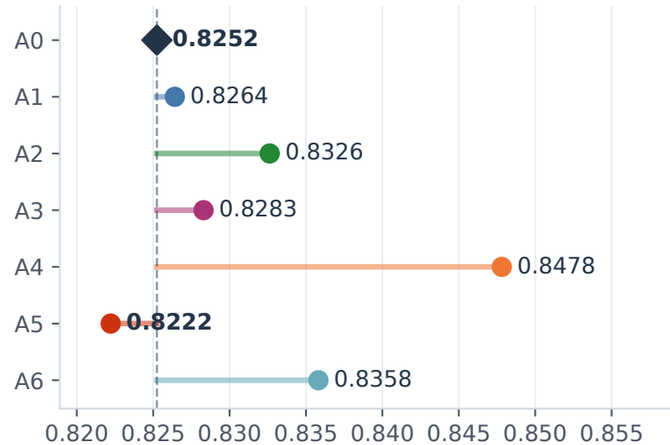

Higher is better

Validation: selected-token CE

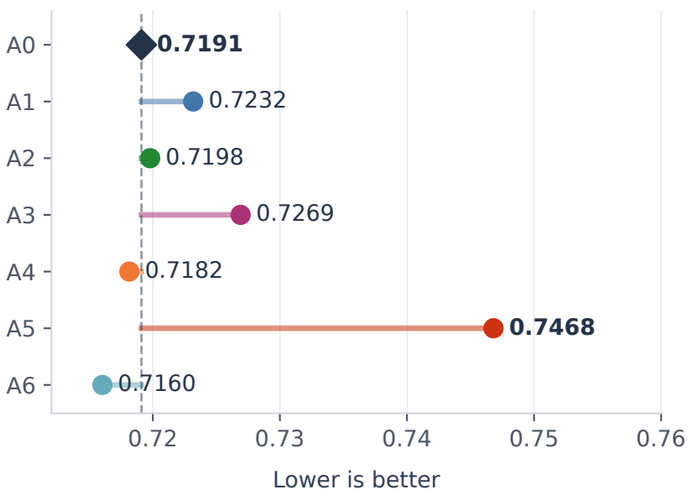

Lower is better

CDR infilling: AAR

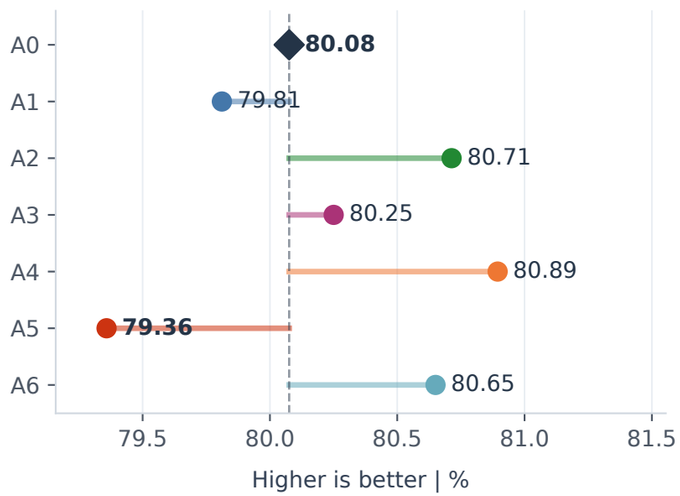

Higher is better | %

Representation: macro AUPRC

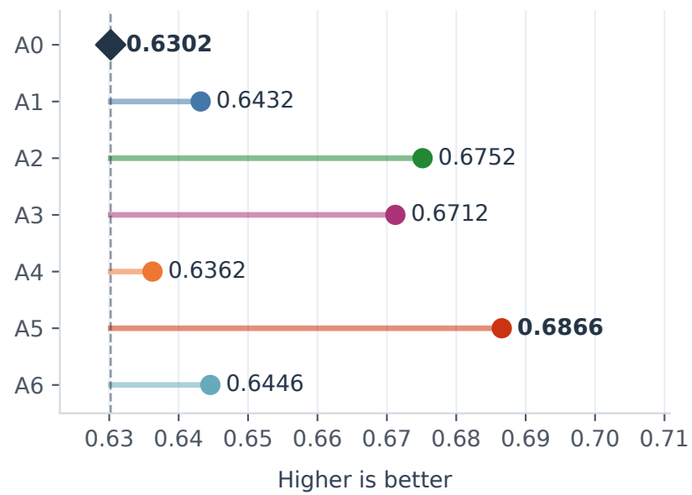

Higher is better

Dashed line = A0 baseline. Each panel uses its own linear scale. Only completed trials appear in this comparison.

Source-macro averages. Generation validity: A0 100.0%, A1 100.0%, A2 91.3%, A3 100.0%, A4 100.0%, A5 100.0%, A6 100.0%.

Exploratory, single seed: no confidence intervals or confirmed winners. Protein sequence similarity is not experimental function.